

UNIVERSITY OF ZAGREB
FACULTY OF ELECTRICAL ENGINEERING AND COMPUTING

MASTER THESIS num. 1417

Deep Learning Model for Base Calling of MinION Nanopore Reads

Marko Ratković

Zagreb, June 2017.

Umjesto ove stranice umetnite izvornik Vašeg rada.
Kako biste uklonili ovu stranicu, obrišite naredbu \izvornik.

Thanks ...

CONTENTS

1. Introduction	1
1.1. Objectives	1
1.2. Organization	1
2. Background	3
2.1. Sequencing	3
2.2. Oxford Nanopore MinION	4
2.2.1. Technology	4
2.3. Existing basecallers	5
2.3.1. Official	5
2.3.2. Third-party	6
3. Methods	8
3.1. Arhitecture	8
3.1.1. RNN	8
3.1.2. CNN	9
3.1.3. Residual Networks	12
3.2. CTC Loss	13
3.2.1. Definition	13
3.2.2. Objective	14
3.2.3. Output decoding	14
3.3. Batch normalization	14
4. Implementation	16
4.1. Deep Learning model	16
4.2. Training	17
4.3. Evaluation methods	19
4.4. Technologies	20

5. Results	22
5.1. Data	22
5.2. Error rates per read	22
5.3. Consensus analysis	26
6. Conclusion	27
Bibliography	28

LIST OF FIGURES

2.1. Depiction of the sequencing process	3
2.2. DNA strain being pulled through a nanopore	5
2.3. Basecalling	7
3.1. An unrolled recurrent neural network	8
3.2. Convolution layer, kernel size 3 with with stride 1.	10
3.3. Activation functions	11
3.4. Dimensionality reduction by pooling, kernel size 2, stride 2	11
3.5. Receptive field after 2 layers of convolutions with kernel size 3	12
3.6. Comparison between classical CNN and CNN with residual connection . .	12
4.1. Used residual block	16
4.2. Learning curve shown in TensorBoard	17
4.3. Dataset preparation	18
4.4. Overview of training pipeline	19
4.5. Consensus from pileup	20
4.6. Overview of evaluation pipeline	20
5.1. Cigar operations histogram over relative position inside read	24
5.2. Cigar operations histogram over relative position inside read	24
5.3. Cigar operations histogram over relative position inside read	25

LIST OF TABLES

5.1. Ecoli R9 basecalled read lengths in base pairs	22
5.2. lambda R9 basecalled read lengths in base pairs	22
5.3. Alignment specifications of Ecoli R9 basecalled reads	23
5.4. Alignment specifications of lambda R9 basecalled reads	23
5.5. Consensus specifications of Ecoli R9 basecalled reads	23
5.6. Consensus specifications of lambda R9 basecalled reads	23

1. Introduction

In recent years, deep learning and usage of deep neural networks have significantly improved the state-of-the-art in many application domains such as computer vision, speech recognition, and natural language processing[1][2]. In this thesis, we present application of deep learning in the fields of Biology and Bioinformatics for analysis of DNA sequencing data.

DNA is a molecule that makes up the genetic material of a cell responsible for carrying the information an organism needs to survive, grow and reproduce. It is a long polymer of simple units called nucleotides attached together to form two long strands that spiral to create a structure called a double helix. The order of these bases is what determines DNA's instructions, or genetic code.

DNA sequencing is the process of determining this sequence of nucleotides. Originally sequencing was very expensive process but during the last couple of decades, the price of sequencing drastically dropped. A significant breakthrough occurred in May 2015 with the release of MinION sequencer by Oxford Nanopore making DNA sequencing inexpensive and available even for small research teams.

Base calling is a process assigning sequence of nucleotides (bases) to the raw data generated by the sequencing device or sequencer. Simply put, it is a process of decoding the raw output from the sequencer.

1.1. Objectives

Goal of this thesis is to show that the accuracy of base calling is dependent on the underlying software and can be improved using machine learning methods. Novel approach for base calling of raw data using convolutional neural networks is presented.

1.2. Organization

Chapter 2 gives more detailed explanation of the problem, background on nanopore sequencing and overview of state-of-the-art basecallers.

Chapter 3 describe used deep learning concepts in detail used later on in later chapters.

Chapter 4 goes into implementation details, training of the deep learning model and explains methods used to evaluate obtained results.

Chapter 5 consists of the results of testing performed on different datasets as well as comparison with state-of-the-art basecallers.

In the end, the Chapter 6 gives a brief conclusion and possible future work and improvements of the developed basecaller.

2. Background

2.1. Sequencing

Sequencing the entire genome of an organism is a difficult problem due to limitations of technology. All sequencing technologies to date have constraints on the length of the strand they can process. These lengths are much smaller than the genome for a majority of organisms, therefore, whole genome shotgun sequencing approach is used. In this approach, multiple copies of the genome are broken up randomly into numerous small fragments that can be processed by the sequencer. Sequenced fragments are called reads.

Genome assembly is the process of reconstructing the original genome from reads and usually starts with finding overlaps between reads. The quality of reconstruction heavily depends on the length and accuracy of the reads produced by the sequencer. Short reads make resolving repetitive regions practically impossible.

Figure 2.1 depicts process of sequencing.

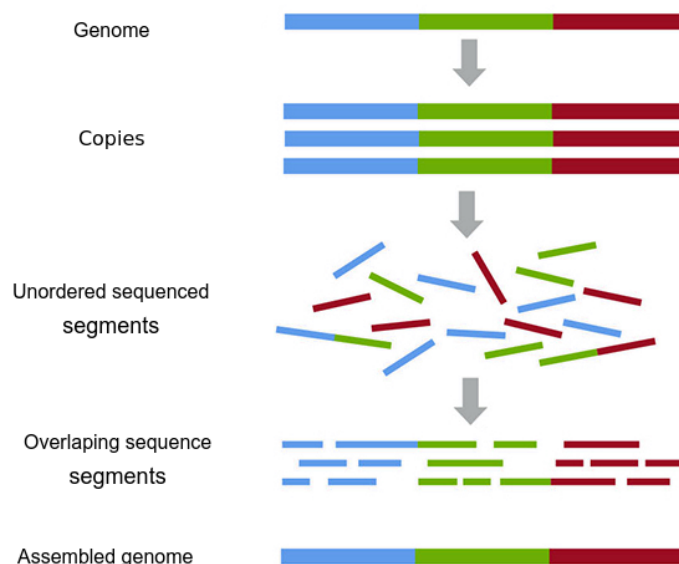


Figure 2.1: Depiction of the sequencing process

Development of sequencing started with work of Frederick Sanger[3][4]. In 1977, he developed the first sequencing method which allowed read lengths up to 1000 bases with

very high accuracy (99.9%) at a cost of 1\$ per 1000 bases[mile_skripta]. Second generation sequencing (IAN Torrent and Illumina devices) reduced the price of sequencing while maintaining high accuracy. Mayor disadvantage of these devices is read length of only a few hundred base pairs. The need for technologies that would produce longer reads led to the development of so-called third generation sequencing technologies. PacBio developed sequencing method that allowed read lengths up to several thousand bases but at a cost of accuracy. Error Rates of PacBio devices are 10-15%.

Cost makes the main obstacle stopping widespread genome sequencing. A significant breakthrough occurred in May 2015 with a release of MinION sequencer by Oxford Nanopore making sequencing drastically less expensive and even portable.

2.2. Oxford Nanopore MinION

The MinION device by Oxford Nanopore Technologies (referenca) is the first portable DNA sequencing device.

It's small weight, low cost, and long read length combined with high-throughput and decent accuracy yield promising results in various applications such as monitoring infectious disease outbreaks [2][3], characterizing structural variants in cancer[4], full human genome assembly [5] what could potentially lead to personalized genomic medicine.

2.2.1. Technology

As its name says, nano-scale pores are used to sequence DNA. An electrical potential is applied over an insulating membrane in which a pore is inserted. As the DNA passes through the pore, the sensor detects changes in ionic current caused by differences in the shifting nucleotide sequences occupying the pore. Figure 2.2 shows change of ionic current as DNA strain being pulled through a nanopore.

Minion offers possibility of sequencing 1D or 2D reads. If information from only one strand is used, the base-calling is termed 1-directional (1D); whereas a 2D base-calling is performed, if information from both strands is incorporated, which results in higher base quality. MinION is able to produce long reads, usually tens of thousand base pairs (with reported reads above 100 thousand pairs[loman]), they have a high sequencing error

Switch from older R7.3 to newer R9 chemistry increases accuracy of produced data. Accuracy of 1D data increased from 70% to 85% while Accuracy of 2D reads from 88% to 94%. video. This increase of accuracy makes 1D reads usable to analysis with benefit of faster sample preparation and faster sequencing than 2D reads.

¹Figure adapted from <https://nanoporetech.com/how-it-works>

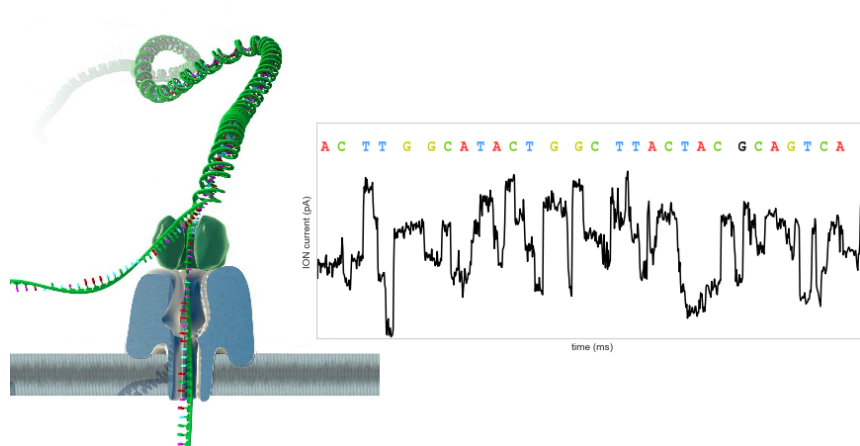


Figure 2.2: DNA strand being pulled through a nanopore ¹

2.3. Existing basecallers

Output of the sequencer is fast5 file. Various fields, including signal. Picture bla shows structure of fast5 file. Analysis fields are added if file is passed to MinKnow for signal segmentation.

2.3.1. Official

Oxford Nanopore has recently, with R9 version of the platform, introduced a variety of base calling options. Some of those are production ready and some experimental. Majority of information regarding differences, specifications and similar is only available through community nanoporetech. <https://community.nanoporetech.com/>

Metrichor, a spin-off of ONT and its main developer of proprietary analysis software, maintains a range of basecallers that have remained the go-to option for most MinION users.

Metrichor is an Oxford Nanopore company, offering cloud-based platform EPI2ME for analysis nanopore data. Even though Metrichor is not a basecaller but rather the platform, its basecaller is often called Metrichor for short because at that point no local basecallers were available. Initially, the Metrichor relied on hidden Markov models (HMM) to find the biological sequence underlying the segmented signal. Preprocess included segmentation of the signal into smaller blocks called events defined by their mean value and the variance of the signal in the block, block length and start location of the block in entire signal.

Metrichor then assumed that each event depends on a context of $k = 6$ consecutive bases and that the context typically shifts by one base in each step. HMM states model context present in the pore and transitions correspond to change of bases in the pore. During transition from one state to another, an event is emitted. Basecalling is performed using Viterbi

algorithm which determines most likely sequence of states for the observed sequence of events. Limitation of this approach is calling homopolymer stretches longer than its k-mer as called bases are determined by difference in two consecutive states.

As of early 2016, with release of R9 chemistry, this model was replaced by a more accurate recurrent neural network (RNN)-implementation. Currently Oxford Nanopore offers several RNN-based local basecaller versions under different names; Albacore, Nanonet and the MinKNOW integrated basecaller. Albacore and the MinKNOW version are stable versions intended for regular MinION users.

Albacore is basecaller by Oxford Nanopore Technologies ready for production and actively supported. The source code of Albacore was not provided and is only available as a binary through the ONT Developer Channel to users who have signed the Developer terms and conditions.

Nanonet² uses the same neural network that is used in Albacore but it is continually under development. As such, it does not contain production code features such as error handling or logging. It uses CURRENNT library for running neural networks.

Scrappie³ is an other basecaller by Oxford Nanopore Technologies. Similar to NanoNet, it is platform for ongoing development. Scrappie was the first basecaller reported to specifically address homopolymer basecalling. It became publicly available just recently in June, 2017.

2.3.2. Third-party

Nanocall[5] was the first third-party open source basecaller for nanopore data. It uses HMM approach like the original R7 Metrichor. Metrichor. It uses the segmented signal from minKNOW and assigns k-mers to the events using a hidden Markov model. Nanocall does not support newer chemistries after R7.3.

DeepNano[6] was the first open-source basecaller based on neural networks that uses bidirectional recurrent neural networks. DeepNano was written in Python, using the Theano library. When released, originally supported R7 chemistry, but support for R9 and R9.4 was added recently.

²<https://github.com/nanoporetech/nanonet/>

³<https://github.com/nanoporetech/scrappie>




Figure 2.3: Basecalling

3. Methods

Process of basecalling can be represented as problem of machine translation where sentence is translated from one language (sequence of events or sequence of current measurements) to another (sequence of nucleotides).

This section explains some key concepts in deep learning needed to understand final model. It gives general idea behind recurrent neural networks and possible problems that serve as motivation for the different approach, usage of convolutional neural networks.

3.1. Arhitecture

3.1.1. RNN

Recurrent neural networks can be viewed as a simple feed-forward network with the twist that the current output does not only depend on the current input but previous inputs as well. RNNs store that information in their hidden state and that state is updated each step. The figure shows simple RNN and the same RNN unfolded in time. Unrolling is simple way of showing how network processes each input in the sequence and updates it's hidden state and is shown in figure 3.1.




Figure 3.1: An unrolled recurrent neural network

These networks are trained using a variant of backpropagation called backpropagation through time which is essentially the same as classical backpropagation on an unfolded network. The gradient is propagated through the entire recurrence relation and the gradient

is multiplied in each step with the same factor, depending on a scale it can make gradient vanish (drop to 0) or exponentially grow each step and explode. These issues are called the vanishing and exploding gradient[7] and are generally resolved by a variant of RNN called *LSTM*[8]. To take into account dependencies of both previous inputs and next inputs in the sequence bidirectional networks (*BRNN*) are used. Idea is to combine two RNN (one in the positive direction, one in negative time direction) and have an output of the current state expressed as a function of hidden states of both RNN and current input. This is the approach used in DeepNano[6].

One of the major drawbacks of this kind of networks is computation time. RNN operate sequentially, the output for the second step depends on the first step and so on, which makes parallelization capabilities of RNN quite limited. This especially goes for Bidirectional RNN.

3.1.2. CNN

Convolutional Neural Networks (CNNs) were responsible for major breakthroughs in Image Classification and are the core of most Computer Vision systems today. More recently CNNs are applied being to problems in Natural Language Processing and have promising results [9][10].

Easiest way to understand a convolution is by thinking of it as a sliding window function applied to a matrix or in case of signal processing vector. The sliding window is called a kernel or a filter. Figure 3.2 shows example of convolution with kernel size 3 and how output is calculated as sum of element-wise multiplication of kernel elements and input vector. Stride defines by how much filter is shifted at each step. By how much the window is moved is step of the convolution and it is called stride. Usually to preserve the same dimension, padding is added to the borders.

Activations

After each convolution layer usually nonlinear layer (or activation layer) is applied. The purpose of this layer is to introduce nonlinearity to a system that basically has just been computing linear operations during the conv layers (just element wise multiplications and summations). In classical neural networks, nonlinear functions like tanh and sigmoid were often used, due to undesirable property of saturation(at either tail of 0 or 1 for sigmoid, -1 or 1 for tanh), other activations are used with these networks. ReLU. The Rectified Linear Unit has become very popular in the last few years. It computes the function $f(x)=\max(0,x)$. In other words, the activation is simply thresholded at zero (see image above on the left). It was found to greatly accelerate the convergence of stochastic gradient descent compared to




Figure 3.2: Convolution layer, kernel size 3 with with stride 1.

the sigmoid/tanh functions [2] and is computationally efficient to calculate.

$$ReLU(x) = \begin{cases} x, & \text{if } x > 0 \\ 0, & \text{otherwise} \end{cases} \quad (3.1)$$

Downside of ReLU saturation to the 0 on one side. Once in this state, neuron is unlikely to recover, because the function gradient at 0 is also 0, so gradient descent learning will not alter the weights. This is the problem known as *dying ReLU*. Different variants of ReLU, PrRelu and ELU are often used to resolve this problem[11][12].

$$PrELU(x) = \begin{cases} x, & \text{if } x > 0 \\ \alpha x, & \text{otherwise} \end{cases} \quad (3.2)$$

$$ELU(x) = \begin{cases} x, & \text{if } x > 0 \\ \alpha(\exp(x) - 1), & \text{otherwise} \end{cases} \quad (3.3)$$

Figure 3.3 show different activation functions.

Pooling

The pooling layer is usually placed after the Convolutional layer. Its primary utility lies in reducing the spatial dimensions of the input for the next Convolutional Layer while keeping the most salient information. Pooling also provides basic invariance to translating (shifting).

Much like the convolution operation performed above, the pooling layer takes a sliding window or a certain size that is moved in stride across the input transforming the values. Usually transformation is performed by taking the maximum value from the values observable in the window (max pooling). Figure 3.4 show dimensionality reduction by pooling by factor of 2 using pool of size 2 with stride 2.




Figure 3.3: Activation functions




Figure 3.4: Dimensionality reduction by pooling, kernel size 2, stride 2

Comparison with RNN

During calculation, each *patch* a convolutional kernel operates on is independent of the other meaning that the entire input layer can be processed concurrently which makes CNNs more efficient than RNNs.

When compared with RNN, in which output can depend on the entire sequence in convolution layer if depends only on limited spatial information in previous layer defined by kernel size. This is called receptive field of the convolution. Figure 3.5 shows each new layers depends on larger portion of the input (z_i sees 5 elements of input). This means that by stacking convolutions sequentially, later layers can see more of the original input.

To calculate final output, input signal has to pass through entire network but as calculations at each layer happen concurrently and each individual computation is small resulting that in practice CNNs have a big speed up over RNNs especially when dealing with large sequences..

As previously show, stacking layers increases receptive field which makes. In the forward pass input flows and transforms, hopefully becoming a representation that is more amenable to our task. During the back phase we propagate a signal, the gradient, back through the network. Just like in RNNs, that signal gets multiplied and depending on the scales it can




Figure 3.5: Receptive field after 2 layers of convolutions with kernel size 3

vanish. Result of that is no gradient flowing to lower layers and no parameter upgrades. On the one hand, we'd like to be able to take in as much context as possible. On the other hand, if we try to increase our receptive fields by stacking layers we risk vanishing gradients and a failure to learn anything. Resnet architecture[13] with its residual layes address this issue.

3.1.3. Residual Networks

A Residual Network or ResNet is a neural network architecture which solves the problem of vanishing gradients using simple trick. Figure 3.6 shows on the left classical CNN that takes input and transforms it using convolution layers and activation. This can be representend as function $H(x)$. $H(x)$ can be written as as sum of some function $F(X)$ and X . $F(X)$ is called the residual. Instead of learning $H(X)$, network learns residual and at the output x is




Figure 3.6: Comparison between classical CNN and CNN with residual connection ¹

simply summed up to the $F(x)$ as shown in the figure.

By stacking these layers, the gradient could theoretically *skip* over all the intermediate layers and reach the bottom without vanishing.

¹Figure adapted from the original paper [13]

3.2. CTC Loss

The goal is to design model which can convert from a sequence of events of current measurements into a sequence of base pairs.

Suppose that we have an input sequence X (signal data) and the desired output sequence Y (nucleotides). X and Y will be of different lengths (the length of base pairs is always smaller than the length of the raw signal), which may pose a problem.

Instead of having a variable size of the output from the neural network, we can limit it to length m and from direct output of the network in some way decode our desired output sequence Y . m is the maximal allowed length of output sequence Y . Idea is that the network outputs is fixed width and the variable length sequence is derived from them. The neural network can be considered to be simply a function that takes in some input sequence X (of length n) and generate some sequence O (of length m). Note that this generated sequence is not same as output sequence Y .

3.2.1. Definition

The key idea behind Connectionist Temporal Classification (CTC) [14] is that instead of directly generating output sequence Y as output from the neural network, we generate a probability distribution at every output length (from $t=1$ to $t=m$) that after *decoding* gives maximum likelihood output sequence Y . Finally, network is trained by creating an objective function that restricts the maximum likelihood decoding for a given sequence X to correspond to our desired target sequence Y .

Given an input sequence X of length n , the network generates some probabilities over all possible labels (A, C, T, and G) with an extra symbol representing a "blank" at each timestep.

The output generated by the network is called *path*. Path is defined by the sequence of its elements $\pi = (\pi_1, \pi_2, \dots, \pi_m)$. The probability of a given path π , given inputs X , can then be written as the product of all its forming elements:

$$P(\pi|X) = \prod_{t=1}^m o_t(\pi_t), \quad (3.4)$$

where $o_t(\pi_t)$ is probability of π_t being t^{th} element on path π

Real output sequence, for given path, is obtained by traversing the path and removing all blanks and duplicate letters. Let $decode(\pi)$ be the output sequence corresponding to a path π . The probability of output sequence Y is then the sum of probabilities of all paths that decode to Y :

$$P(Y|X) = \sum_{\pi \in decode^{-1}(Y)} P(\pi|X) \quad (3.5)$$

3.2.2. Objective

Given the dataset $D = \{(X, Y)\}$, training objective is the maximization of the likelihood of each training sample which corresponds to the minimization of negative log likelihood:

$$L(D) = - \sum_{(X,Y) \in D} \ln P(Y|X) \quad (3.6)$$

3.2.3. Output decoding

Given the probability distribution $P(Y|X)$ and given input sequence X , most likely Y^* can be computed.

$$Y^* = \operatorname{argmax}_{Y \in L^m} P(Y|X) = \operatorname{argmax}_{Y \in L^m} \sum_{\pi \in \text{decode}^{-1}(Y)} P(\pi|X), \quad (3.7)$$

where L^m set of all possible sequences over alphabet L

with length less than or equals to m

The probability of a single output sequence Y is the sum of probabilities of all paths that decode to Y and the most probable sequence is needed to be found. Calculation of all possible sequences is computationally intractable but exist algorithms that approximate decoding.

One naive possibility is to take the most probable path and say that output sequence corresponds to that path. This is not necessarily true: suppose we have one path with probability 0.1 corresponding with sequence A , and ten paths with probabilities 0.05 each corresponding to sequence B . Clearly, label B is preferable overall, since it has an overall probability of 0.5; however, this naive best path decoding would select label A , which has a higher probability than any single path for label B .

Better approximations can be calculated using beam search algorithm proposed in paper[15]. Idea behind this approach is to

This serves as a brief overview of the method and explains key concepts why it is used. More detailed explanation can be found in original paper[14] or various blog post[16].

3.3. Batch normalization

Batch normalization is method proposed in paper[17] that accelerates learning process. During training parameters are updated and distribution of outputs of each layers keep changing. Small change in distribution of outputs in early layers can cause drastic change in later layers and those layers need to adapt to the new scale of their inputs. This change of distribution is called the internal covariate shift and results in slows learning. Proposed solution for this

is to center each output from activations of the mini-batch to zero-mean and unit variance. After that learned scale and offset are applied. This process is called batch normalization. After training, mean and variance for each activation is computed on the whole training dataset rather than on mini-batches during training. Pros of using batch normalization Batch normalization offers several advantages other than reducing internal covariant shift. It offers more robust learning process by reducing dependance on scale of the parameters and their initial values. This allows usage of larger learning rates and faster learning all together. It is show in the original paper that batch normalisation also regularizes the model that could potentially improve performance of the model.

4. Implementation

4.1. Deep Learning model

Final model is residual neural network consisting of 72 residual blocks depicted in figure 4.1. This is a variant of architecture proposed in paper [18] with the difference of ELU being used as activation instead of ReLU as it is reported[19] to speeds up learning process and improve accuracy as the depth increase.

It is difficult to discuss actual performance benefit over simple ReLU without more detailed analysis and comparison for this problem.

Each residual block contains 2 convolution layers making total number of convolution layers 144. Each convolutional layer in this models uses 64 kernels of 3. Pooling with kernel size 2 is used every 48 layers making output 8 times shorter than signal input. This is used to help training by reducing number of required blank labels in the output and reducing computation effort.




Figure 4.1: Used residual block

4.2. Training

To train described model we need to obtain dataset that consist of sample pairs (X_i, Y_i) where (X_i) is input signal and Y_i is output sequence. One of the mayor issues is determining correct output for given signal. One option is to use existing basecaller like Metrichor to determine output sequences.




Figure 4.2: Learning curve shown in TensorBoard

Supervised learning so for need to specify dataset that for each input signal we need specify desired sequence of bases. X_i, Y_i where len of X and len of Y_i are not specified.

For each input file, ground truth is not specified. If we use output of Metrichor or any other basecaller we limit our model to obtain accuracy of used basecaller in best case. We limit our train data to only sequencing data of know organisms (organisms with know reference genome) and try to correct data by aligning the read produced by metrichor or any other basecaller to reference genome. Alignment destination is used as target sequencing in training. Using Metrichor basecalled data we can determine for each called event values such as start in signal, length of event, k-mer state in the pore and using move field change of k-mer from previous state.

Used dataset for training consists of raw signal from fast5 files split raw data into blocks of size l . For each block it is easy to determine basecalled events from Metrichor that belong to particular block using start information and from those events basecalled sequence.

$$block_index = \frac{event_start * sampling_rate}{block_size}$$

Full basecalled sequence is aligned to the reference genome and alignment was obtained. For each block target sequence is determined from alignment information.

Figure X shows how for given signal block

Figure bla shows augmentation of data.

Gradient descent is one of the most popular algorithms to perform optimization and by far

the most common way to optimize neural networks. Gradient descent is a way to minimize an objective function parameterized by a models parameters by updating the parameters in the opposite direction of the gradient of the objective function to the parameters.

Adam is a stochastic gradient descent algorithm based on estimation of 1st and 2nd-order moments. It is often used as it offers fast and stable convergenton, even using default arguments proposed by authors. Learning rate 1e-3 with exponential decay 0.98 every 10000 steps, batch size was set to 8.

To eliminate possibility of overfitting to the know reference, model is trained and tested on reads that align to different regions of reference genome and those regions should not overlap. Also test is conducted on separate set of sequencing data for different organism than the one used for training.




Figure 4.3: Dataset preparation




Figure 4.4: Overview of training pipeline

4.3. Evaluation methods

To evaluate trained model we base call test set, align those reads to reference and calculate various statistics on align data. From cigar string it is easy to calculate following: the proportion of bases in a sequencing ‘read’ that align to a matching base in a reference sequence

$$read_len = n_matches + n_mismatches + n_insertions \quad (4.1)$$

$$match_rate = \frac{n_matches}{read_length} \quad (4.2)$$

$$mismatch_rate = \frac{n_mismatches}{read_length} \quad (4.3)$$

$$insertion_rate = \frac{n_insertions}{read_length} \quad (4.4)$$

$$deletion_rate = \frac{n_deletion}{read_length} \quad (4.5)$$

$$match_rate = \frac{n_matches}{read_length} \quad (4.6)$$

Results are calculated for each read in test dataset and median value, mean and standard deviation is expressed for the whole dataset.

To validate consistency of a basecaller, basecalled data is aligned to the reference genome and consensus sequence is called from all reads covering single position. Consensus sequence is compared with the reference genome and following measures are calculated:

$$identity_percentage = 100 * \frac{n_correct_bases}{reference_length} \quad (4.7)$$

$$match_rate = \frac{n_correct_bases}{consensus_length} \quad (4.8)$$

$$snp_rate = \frac{n_snp}{consensus_length} \quad (4.9)$$

$$insertion_rate = \frac{n_insertions}{consensus_length} \quad (4.10)$$

$$deletion_rate = \frac{n_deletion}{consensus_length} \quad (4.11)$$




Figure 4.5: Consensus from pileup




Figure 4.6: Overview of evaluation pipeline

4.4. Technologies

Overall solution was implemented in Python programming language. Described model is implemented using TensorFlow. It is an open source software library for numerical computation using data flow graphs developed by Google. TensorFlow, even tho is considered low-level

framework offers implementations of higher level concepts (layers, losses, and optimizers) which makes it great for prototyping while keeping it modular and extensible for highly specific tasks as well.

TensorFlow offers efficient GPU implementations of various layers and losses but as of version 1.2 lacks GPU implementation of used CTC loss, so WARP-CTC¹ was used. It offers both GPU and CPU implementations as well as bindings for TensorFlow.

For alignment tasks, developed tool offers support for GraphMap and BWA but can easily be extended with any other aligner that outputs results in Sam file format.

SAMTools² and it's Python bindings PySam³ were used for conversions between various file formats used in Bioinformatics.

Docker was used for automating the deployments on different machines. It helps us resolve problem know as *dependency hell*⁴ keeping all dependencies in single container thus eliminating possible conflict between packages on host OS. Nvidia Docker⁵ was used for GPU support.

All training was done on the server with *Intel(R) Xeon(R) E5-2640 CPU*, 600 GB of RAM and *NVIDIA TITAN X Black* with 6GB of GDDR5 memory and 2880 CUDA cores.

¹<https://github.com/baidu-research/warp-ctc>

²<http://www.htslib.org/>

³<https://github.com/pysam-developers/pysam>

⁴https://en.wikipedia.org/wiki/Dependency_hell

⁵<https://github.com/NVIDIA/nvidia-docker>

5. Results

5.1. Data

5.2. Error rates per read

Table 5.1: Ecoli R9 basecalled read lengths in base pairs

	median	mean	std
deepnano	5526.5	8126.694000	7406.554786
metrichor	5809.5	8933.275000	9189.709720
nanonet	3286.5	4874.406582	4803.182344
resdeep	5784.0	8990.988989	9297.972688

Table 5.2: lambda R9 basecalled read lengths in base pairs

	median	mean	std
deepnano	4740.0	4664.750000	2628.512543
metrichor	5491.0	5482.952941	2748.446253
nanonet	4931.5	4925.804878	2739.987512
resdeep	5229.0	5138.764706	2605.958080

Figure 5.1 shows how cigar operations are distributed across the reads. It is shown that mismatches and insertion occur more frequently at the ends. This is not only the case for mincall, other basecallers show same property. This may be caused by incorrect detection of read start and end.

Table 5.3: Alignment specifications of Ecoli R9 basecalled reads

	Match % (median)	Mismatch % (median)	Insertion % (median)	Deletion % (median)
deepnano	90.254762	6.452852	3.274420	11.829965
metrichor	90.560455	5.688105	3.660381	8.328271
nanonet	90.607674	5.608912	3.652791	8.299046
resdeep	91.408591	5.019141	3.477739	7.471608

Table 5.4: Alignment specifications of lambda R9 basecalled reads

	Match % (median)	Mismatch % (median)	Insertion % (median)	Deletion % (median)
deepnano	86.997687	9.623494	3.442490	16.052830
metrichor	87.714988	7.835052	4.093851	10.757491
nanonet	88.415611	8.178372	3.629653	11.793022
resdeep	89.694482	7.238095	3.078796	13.450292

Table 5.5: Consensus specifications of Ecoli R9 basecalled reads

	Match %	Snp %	Insertion %	Deletion %
deepnano	98.874222	1.004407	0.121371	0.904092
metrichor	99.122300	0.746359	0.131342	0.629992
nanonet	97.969082	1.570034	0.460885	1.515800
resdeep	99.236079	0.647438	0.116482	0.550985

Table 5.6: Consensus specifications of lambda R9 basecalled reads

	Match %	Snp %	Insertion %	Deletion %
deepnano	99.344256	0.643333	0.012412	0.264780
metrichor	99.562607	0.418824	0.018569	0.146485
nanonet	99.442586	0.540898	0.016516	0.196127
resdeep	99.541180	0.440219	0.018601	0.297613




Figure 5.1: Cigar operations histogram over relative position inside read

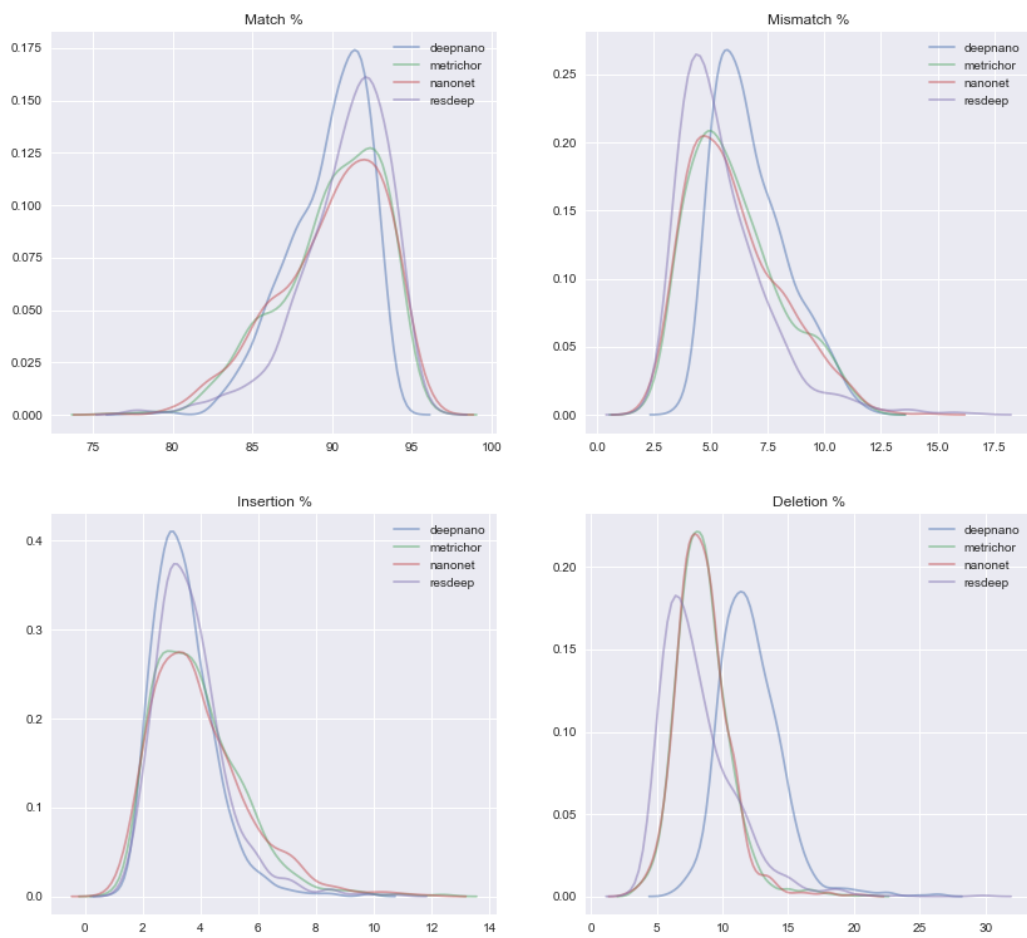


Figure 5.2: Cigar operations histogram over relative position inside read

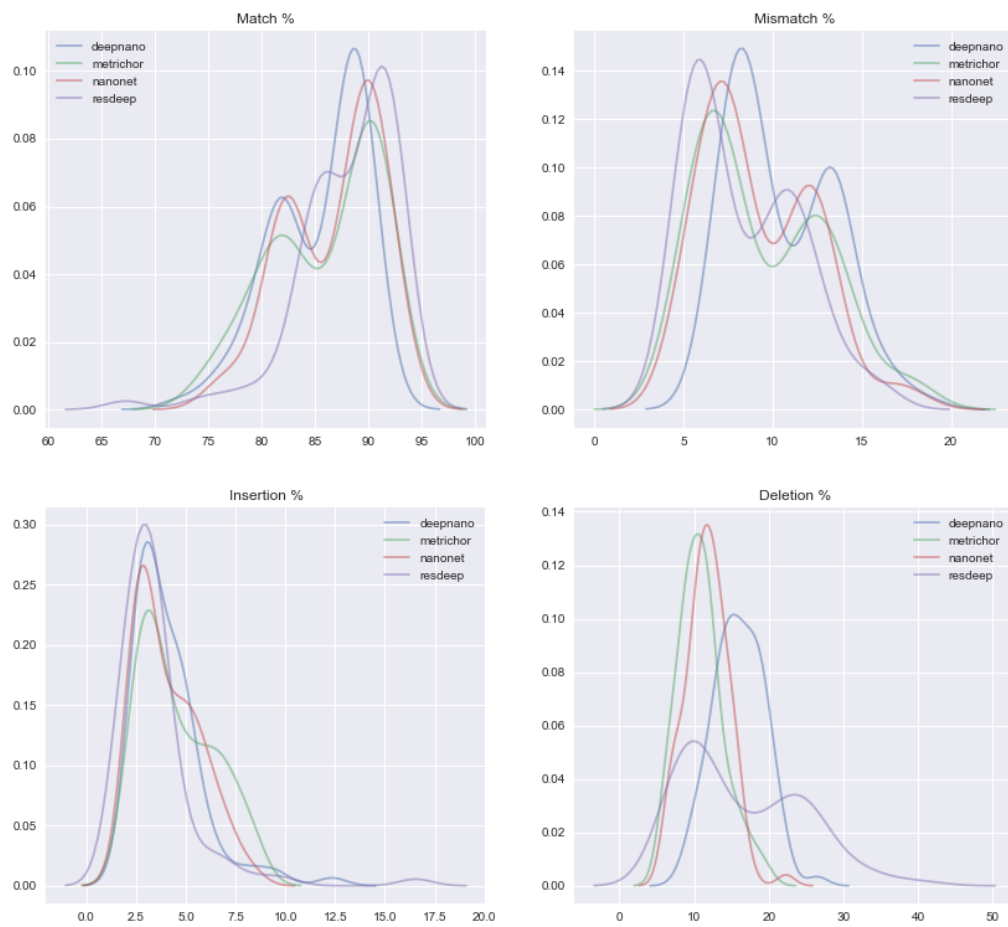


Figure 5.3: Cigar operations histogram over relative position inside read

5.3. Consensus analysis

6. Conclusion

Conclusion.

BIBLIOGRAPHY

- [1] Yann LeCun and Yoshua Bengio. The handbook of brain theory and neural networks. chapter Convolutional Networks for Images, Speech, and Time Series, pages 255–258. MIT Press, Cambridge, MA, USA, 1998. ISBN 0-262-51102-9. URL <http://dl.acm.org/citation.cfm?id=303568.303704>.
- [2] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 25*, pages 1097–1105. Curran Associates, Inc., 2012. URL <https://goo.gl/UpFBv8>.
- [3] Mirjana Domazet-Lošo Mile Šikić. *Bioinformatika*. Bioinformatics - course matherials, Faculty of Electrical Engineering and Computing, University of Zagreb, 2013.
- [4] Erik Pettersson, Joakim Lundeberg, and Afshin Ahmadian. Generations of sequencing technologies. *Genomics*, 93(2):105–111, feb 2009. doi: 10.1016/j.ygeno.2008.10.003. URL <https://doi.org/10.1016/j.ygeno.2008.10.003>.
- [5] Matei David, Lewis Jonathan Dursi, Delia Yao, Paul C Boutros, and Jared T Simpson. Nanocall: An open source basecaller for oxford nanopore sequencing data. *bioRxiv*, 2016. doi: 10.1101/046086. URL <http://biorxiv.org/content/early/2016/03/28/046086>.
- [6] Vladimír Boža, Broňa Brejová, and Tomáš Vinař. DeepNano: Deep recurrent neural networks for base calling in MinION nanopore reads. *PLOS ONE*, 12(6):e0178751, jun 2017. doi: 10.1371/journal.pone.0178751. URL <https://doi.org/10.1371/journal.pone.0178751>.
- [7] Denny Britz. *Recurrent Neural Networks Tutorial - Backpropagation Through Time and Vanishing Gradients*. URL <https://goo.gl/Qkv9gC>. [Online; posted 15-September-2015].
- [8] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.

- [9] Nal Kalchbrenner, Lasse Espeholt, Karen Simonyan, Aaron van den Oord, Alex Graves, and Koray Kavukcuoglu. Neural machine translation in linear time, 2016.
- [10] Jonas Gehring, Michael Auli, David Grangier, Denis Yarats, and Yann N. Dauphin. Convolutional sequence to sequence learning, 2017.
- [11] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification, 2015.
- [12] Djork-Arné Clevert, Thomas Unterthiner, and Sepp Hochreiter. Fast and accurate deep network learning by exponential linear units (elus), 2015.
- [13] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition, 2015.
- [14] Alex Graves, Santiago Fernández, Faustino Gomez, and Jürgen Schmidhuber. Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks. In *Proceedings of the 23rd International Conference on Machine Learning*, ICML '06, pages 369–376, New York, NY, USA, 2006. ACM. ISBN 1-59593-383-2. doi: 10.1145/1143844.1143891. URL <http://doi.acm.org/10.1145/1143844.1143891>.
- [15] Alex Graves and Navdeep Jaitly. Towards end-to-end speech recognition with recurrent neural networks. In Eric P. Xing and Tony Jebara, editors, *Proceedings of the 31st International Conference on Machine Learning*, volume 32 of *Proceedings of Machine Learning Research*, pages 1764–1772, Beijing, China, 22–24 Jun 2014. PMLR. URL <http://proceedings.mlr.press/v32/graves14.html>.
- [16] Andrew Gibiansky. *Speech Recognition with Neural Networks*. URL <http://andrew.gibiansky.com/blog/machine-learning/speech-recognition-neural-networks/>. [Online; posted 23-April-2014].
- [17] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift, 2015.
- [18] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Identity mappings in deep residual networks, 2016.
- [19] Anish Shah, Eashan Kadam, Hena Shah, Sameer Shinde, and Sandip Shingade. Deep residual networks with exponential linear unit. 2016. doi: 10.1145/2983402.2983406.

Deep Learning Model for Base Calling of MinION Nanopore Reads

Abstract

Abstract.

Keywords: Keywords.

Model dubokog učenja za određivanje očitanih baza dobivenih uređajem za sekvenciranje MinION

Sažetak

Sažetak na hrvatskom jeziku.

Ključne riječi: Ključne riječi, odvojene zarezima.